NWI-I00041
Information Retrieval – Lecture 4

Dr. ir. Faegheh Hasibi
f.hasibi@cs.ru.nl

Prof.dr.ir. Arjen P. de Vries
arjen@acm.org

Nijmegen, September 21st, 2020

Abstract IR Architecture

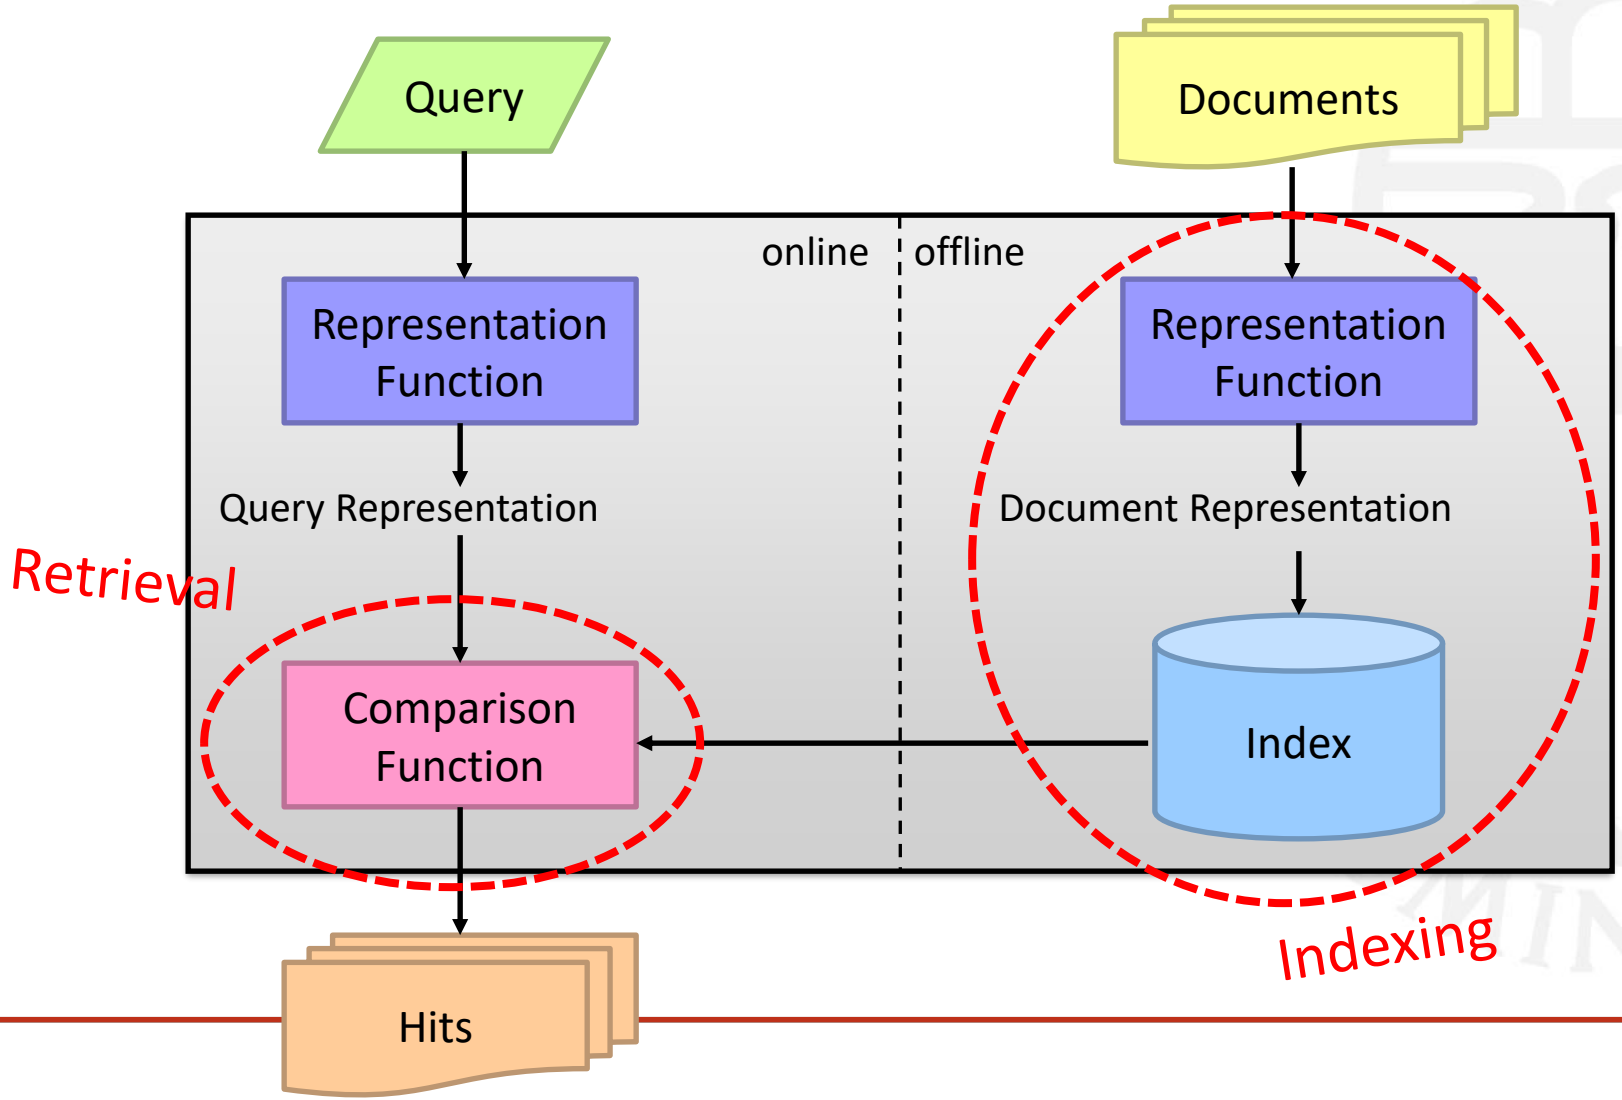

Inverted Index

- Indexes are data structures designed to make search faster
- The inverted index is an efficient & flexible data structure
- Default technique to rank documents efficiently on their similarity to the query
- Find all exact matches efficiently
  - Assumes the query is “correct”

Inverted Index

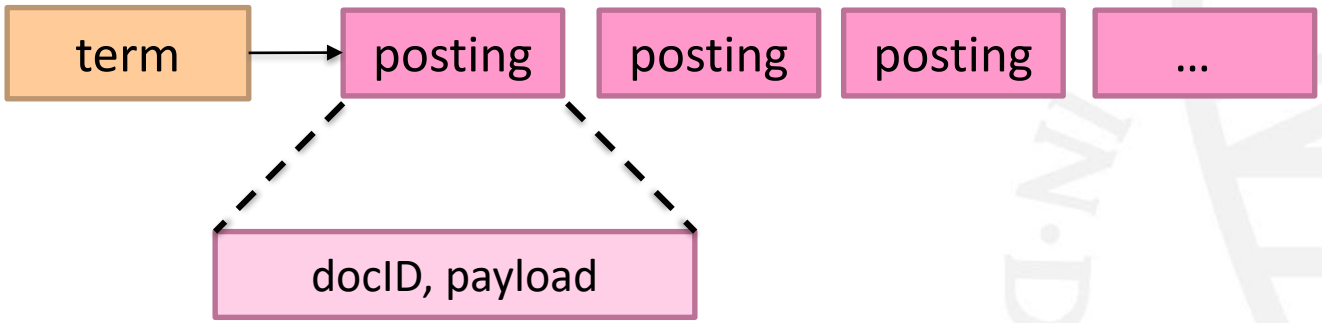

Payload (optional): other associated information such as count and position

Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

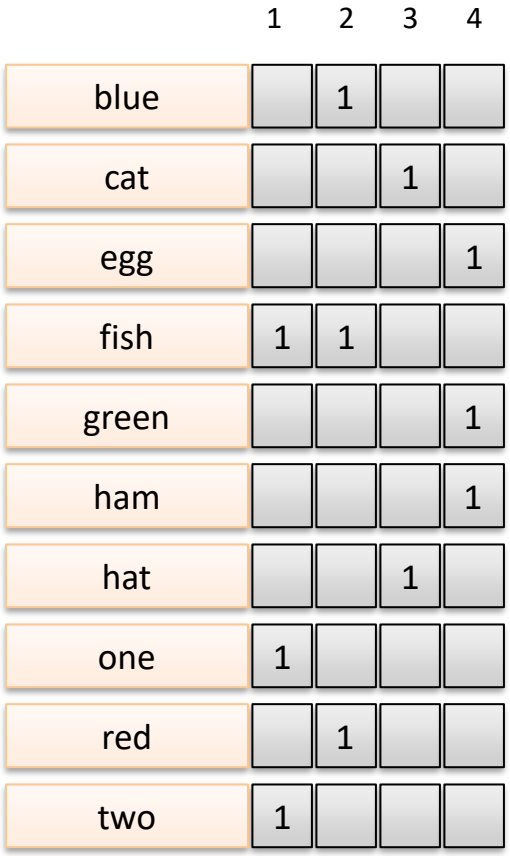

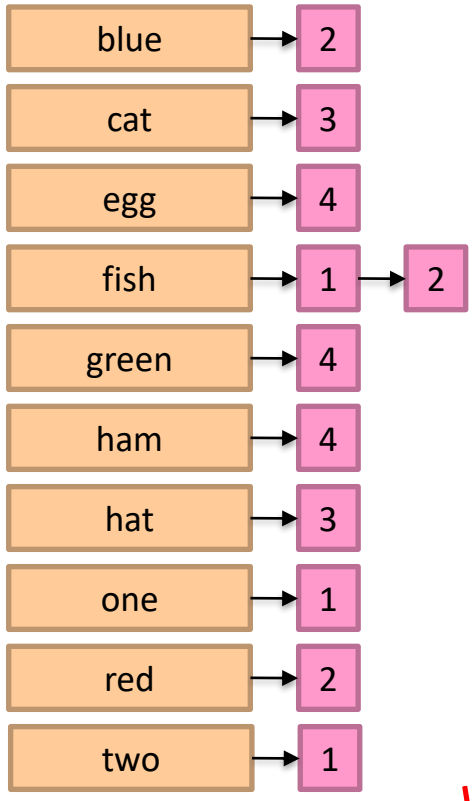

postings lists
(often in sorted order)

Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

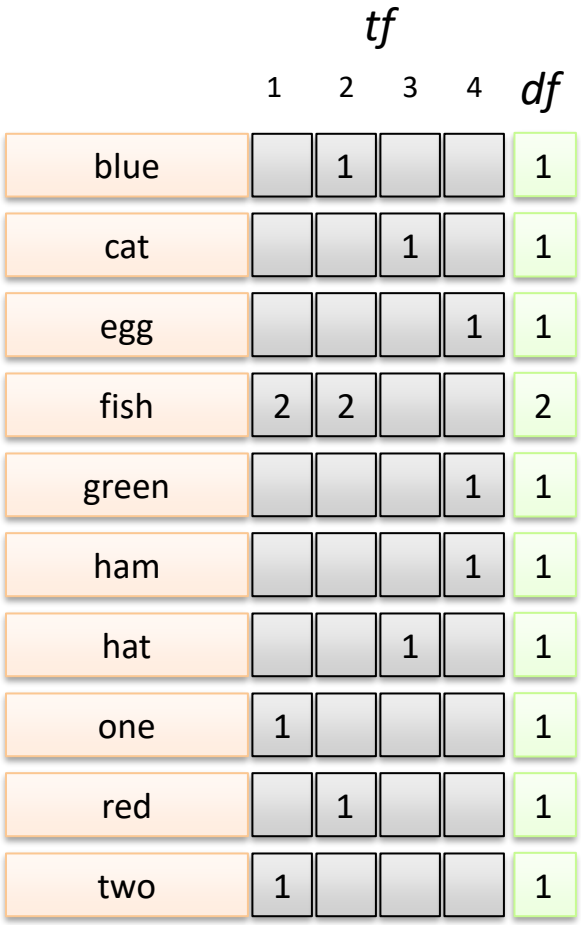

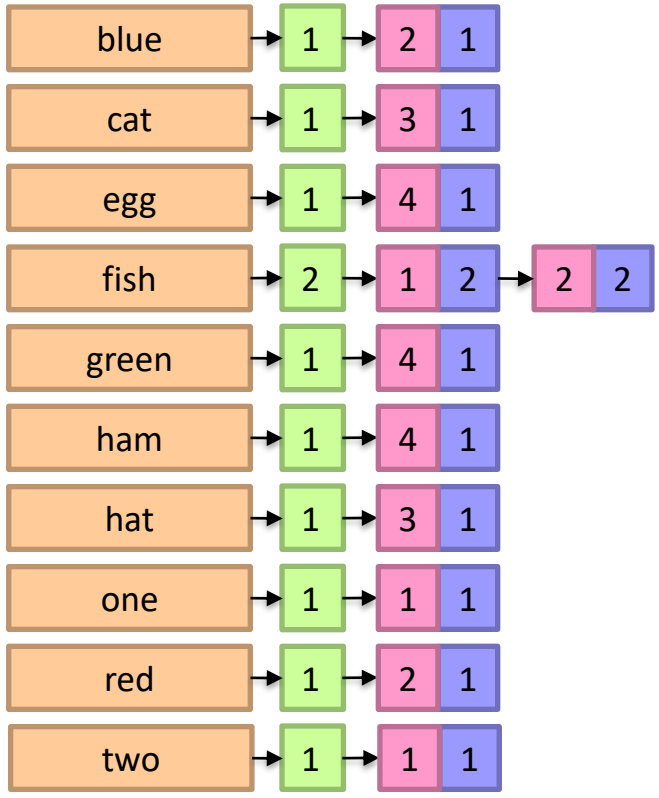

Doc 1 one fish, two fish

Doc 2 red fish, blue fish

Doc 3 cat in the hat

Doc 4 green eggs and ham

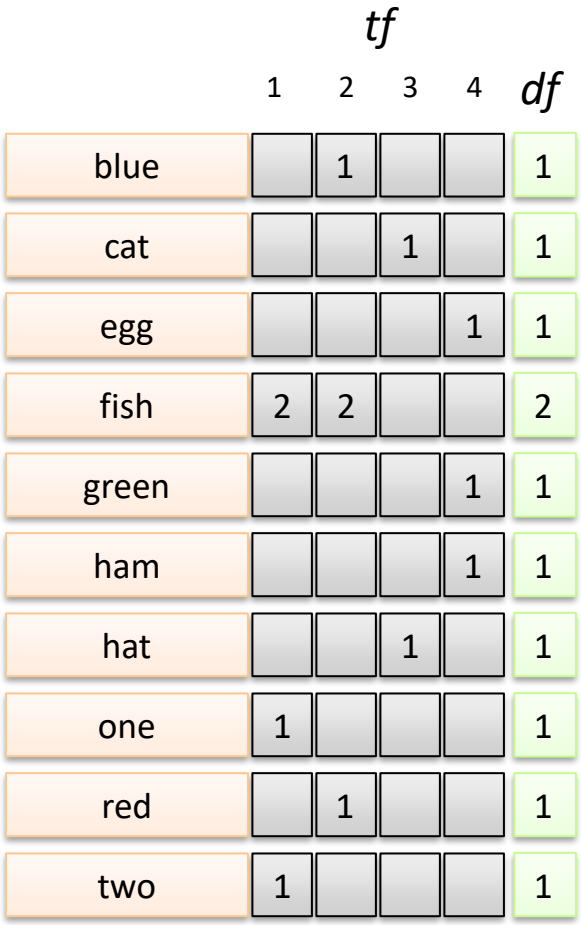

Compression

- Inverted lists are very large, stored on disk
- Compression of the postings lists reduces the bandwidth needed to read from disk
- Compression algorithms optimized for high decompression speed (~3 GB/s and up; PFORDelta invented at CWI Amsterdam is widely used)

Compression Methods

- Bit-aligned
  - Unary codes
  - γ/δ codes
  - Golomb codes (local Bernoulli model)

- Byte-aligned technique
  - Vbyte

VByte

- Simple idea: use only as many bytes as needed
  - Need to reserve one bit per byte as the “continuation bit”
  - Use remaining bits for encoding value

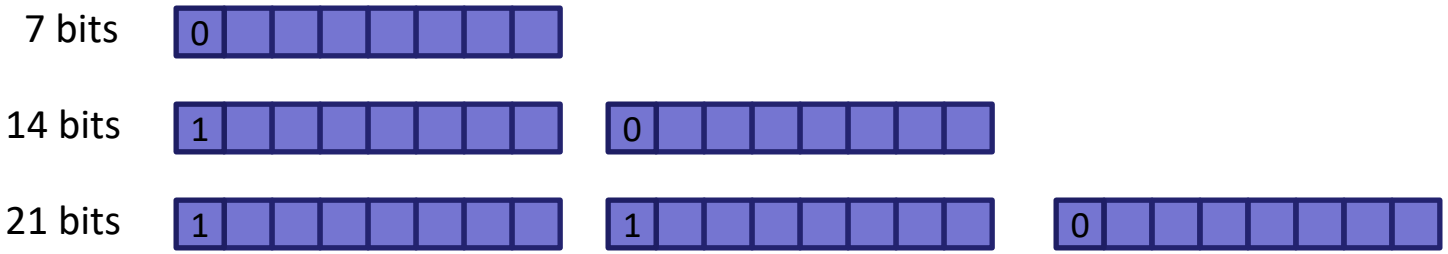

What to compress?!

- Gap-encoding: compress the numbers that result from the difference between consecutive document IDs
  - Frequent terms: small numbers – better compression as a result
  - Infrequent terms: still large numbers, likely unique; but very few relatively, given Zipf's law

Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

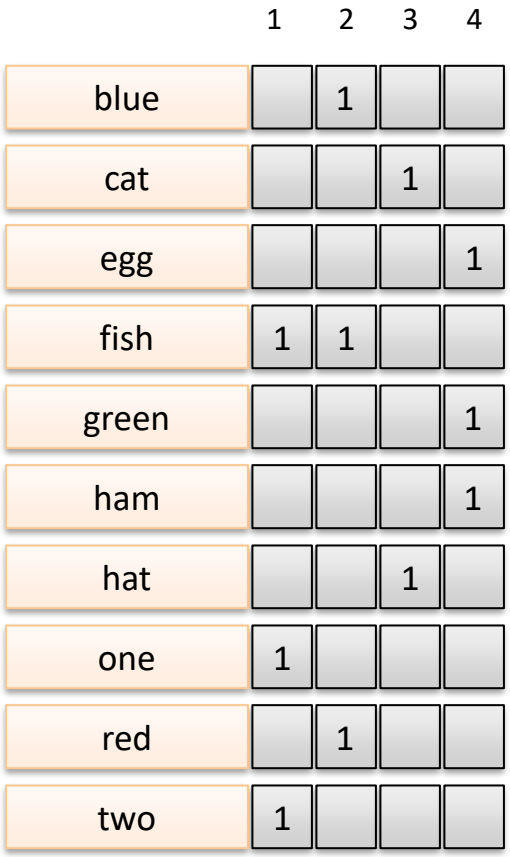

Indexing: building this structure

Retrieval: manipulating this structure

Retrieval in a Nutshell

- Look up postings lists corresponding to query terms
- Traverse postings for each query term
- Store partial query-document scores in accumulators
- Select top *k* results to return

Retrieval and Query Processing

- Strategies for processing the data in the index for producing query results

- Document-at-a-time
  - Calculates complete scores for documents by processing all term lists, one document at a time

- Term-at-a-time
  - Accumulates scores for documents by processing term lists one at a time

Both approaches have optimization techniques that significantly reduce time required to generate scores

Document-at-a-Time

- Evaluate documents one at a time (score all query terms)

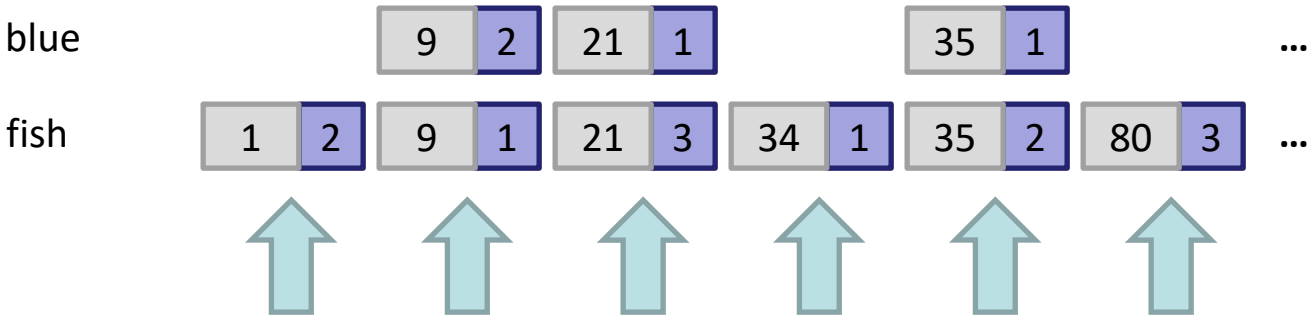

Accumulators
(e.g. min heap)

Document score in top k?

Yes: Insert document score, extract-min if heap too large

No: Do nothing

- Trade-offs
  - Small memory footprint (good)
  - Skipping possible to avoid reading all postings (good)
  -

Term-At-A-Time

- Evaluate documents one query term at a time
  - Usually, starting from most rare

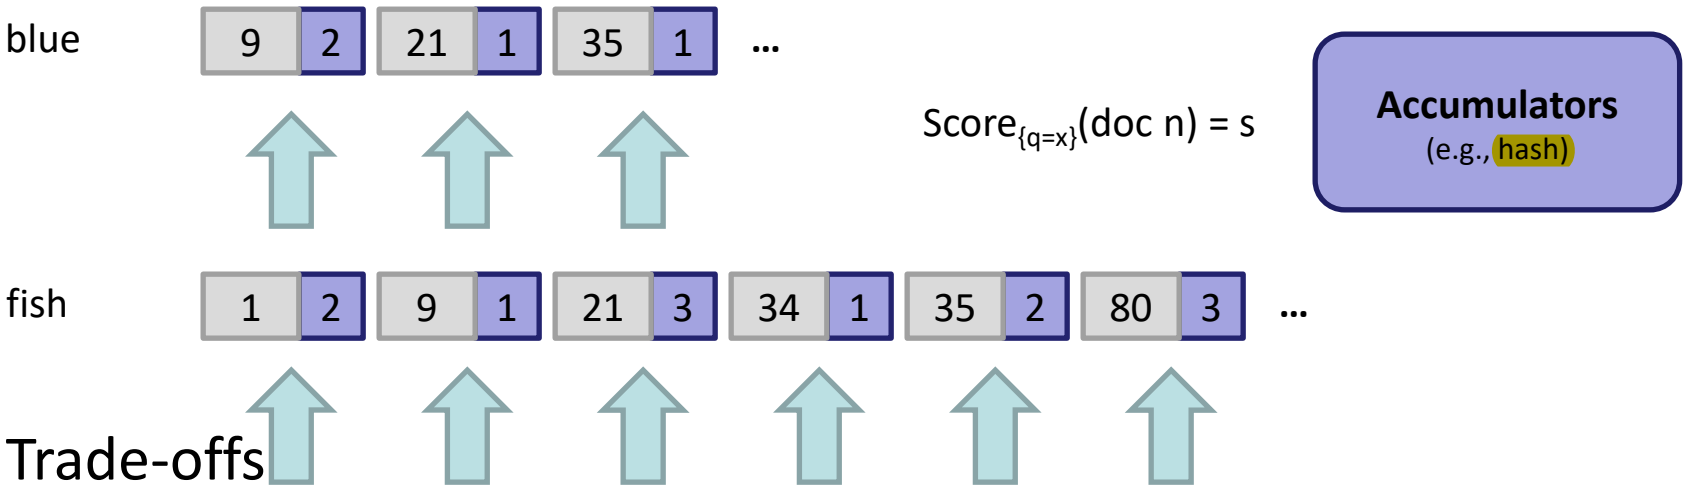

- Trade-offs
  - Early termination heuristics (good)
  - Large memory footprint (bad), but filtering heuristics possible